

```
# 1.Opracować przepływ pracy uczenia maszynowego zagadnienia
klasyfikacji (pojedyncze drzewo decyzyjne)
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, export_text
from sklearn.metrics import classification_report, accuracy_score

data = pd.read_csv('smokers.csv')
data.dropna()

data['Smoking Status'] = data['Smoking Status'].map({'current': 1,
'former': 0, 'never': -1})
data['Gender'] = data['Gender'].map({'f': 0, 'm': 1})

X = data.drop(columns=['GSM', 'Smoking Status'])
y = data['Smoking Status']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

clf = DecisionTreeClassifier(random_state=42)
clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test,
y_pred))
print("\nDecision Tree:\n")
print(export_text(clf, feature_names=list(X.columns)))
```

Accuracy: 0.656934306569343

Classification Report:

	precision	recall	f1-score	support
-1	0.39	0.36	0.37	39
1	0.75	0.78	0.76	98
accuracy			0.66	137
macro avg	0.57	0.57	0.57	137
weighted avg	0.65	0.66	0.65	137

Decision Tree:

```
| --- cg02839557 <= 0.04
|   | --- Age <= 32.00
|   |   | --- class: -1
|   |   | --- Age > 32.00
|   |   | --- cg00213748 <= 0.92
|   |   |   | --- cg03052502 <= 0.99
```

```

| | | |--- class: 1
| | | |--- cg03052502 > 0.99
| | | |--- class: -1
| | | |--- cg00213748 > 0.92
| | | |--- class: -1
--- cg02839557 > 0.04
| | | |--- cg01707559 <= 0.06
| | | |--- cg03695421 <= 0.58
| | | |--- class: 1
| | | |--- cg03695421 > 0.58
| | | |--- class: -1
--- cg01707559 > 0.06
| | | |--- cg00050873 <= 0.57
| | | |--- cg01707559 <= 0.22
| | | |--- cg03052502 <= 0.45
| | | |--- Age <= 62.50
| | | |--- cg03443143 <= 0.38
| | | |--- cg02004872 <= 0.26
| | | |--- class: -1
| | | |--- cg02004872 > 0.26
| | | |--- class: 1
| | | |--- cg03443143 > 0.38
| | | |--- Age <= 48.50
| | | |--- cg02842889 <= 0.39
| | | |--- cg01707559 <= 0.22
| | | |--- class: 1
| | | |--- cg01707559 > 0.22
| | | |--- class: -1
| | | |--- cg02842889 > 0.39
| | | |--- class: -1
| | | |--- Age > 48.50
| | | |--- cg02004872 <= 0.10
| | | |--- class: -1
| | | |--- cg02004872 > 0.10
| | | |--- class: 1
| | | |--- Age > 62.50
| | | |--- cg02494853 <= 0.08
| | | |--- class: -1
| | | |--- cg02494853 > 0.08
| | | |--- class: 1
--- cg03052502 > 0.45
| | | |--- cg03155755 <= 0.45
| | | |--- class: 1
| | | |--- cg03155755 > 0.45
| | | |--- cg03683899 <= 0.33
| | | |--- class: 1
| | | |--- cg03683899 > 0.33
| | | |--- class: -1
--- cg01707559 > 0.22

```



```

--- cg03052502 > 0.52
    |--- cg03244189 <= 0.35
    |   |--- class: 1
    |   |--- cg03244189 > 0.35
    |   |--- class: -1
--- cg00455876 > 0.35
    |--- cg00455876 <= 0.37
    |   |--- cg02842889 <= 0.40
    |   |   |--- cg03695421 <= 0.21
    |   |   |   |--- class: -1
    |   |   |--- cg03695421 > 0.21
    |   |   |   |--- class: 1
    |   |--- cg02842889 > 0.40
    |   |   |--- class: -1
--- cg00455876 > 0.37
    |--- cg02494853 <= 0.09
    |   |--- cg00212031 <= 0.58
    |   |   |--- cg03706273 <= 0.13
    |   |   |   |--- cg02494853 <= 0.06
    |   |   |   |   |--- cg02842889 <= 0.33
    |   |   |   |   |   |--- class: -1
    |   |   |   |   |   |--- cg02842889 > 0.33
    |   |   |   |   |   |--- truncated branch of

```

```

|--- cg02494853 > 0.06
|--- class: 1
|--- cg03706273 > 0.13
|--- cg02842889 <= 0.49
|--- class: -1
|--- cg02842889 > 0.49
|--- class: 1
|--- cg00212031 > 0.58
|--- class: -1
|--- cg02494853 > 0.09
|--- class: -1
|--- cg00050873 > 0.57
|--- cg02233190 <= 0.02
|--- class: -1
|--- cg02233190 > 0.02
|--- cg02494853 <= 0.11
|--- Age <= 48.50
|--- cg03695421 <= 0.64
|--- cg02494853 <= 0.05
|--- cg03052502 <= 0.98
|--- cg02494853 <= 0.03
|--- class: -1
|--- cg02494853 > 0.03
|--- class: 1

```



```

X_train = imputer.fit_transform(X_train)
X_test = imputer.transform(X_test)

models = {
    "Random Forest": RandomForestClassifier(n_estimators=100,
random_state=42),
    "Bagging": BaggingClassifier(n_estimators=100, random_state=42),
    "Boosting (AdaBoost)": AdaBoostClassifier(n_estimators=100,
random_state=42),
    "Gradient Boosting": GradientBoostingClassifier(n_estimators=100,
random_state=42)
}

```

```

for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    print(f"=== {name} ===")
    print(f"Accuracy: {accuracy_score(y_test, y_pred):.2f}")
    print(classification_report(y_test, y_pred))

```

=== Random Forest ===

Accuracy: 0.74

	precision	recall	f1-score	support
-1	0.50	0.21	0.29	53
1	0.77	0.93	0.84	152
accuracy			0.74	205
macro avg	0.64	0.57	0.57	205
weighted avg	0.70	0.74	0.70	205

```

c:\Programs\Python\3.11\Lib\site-packages\sklearn\impute\
_base.py:635: UserWarning: Skipping features without any observed
values: ['Gender']. At least one non-missing value is needed for
imputation with strategy='most_frequent'.

```

```
warnings.warn(
```

```

c:\Programs\Python\3.11\Lib\site-packages\sklearn\impute\
_base.py:635: UserWarning: Skipping features without any observed
values: ['Gender']. At least one non-missing value is needed for
imputation with strategy='most_frequent'.

```

```
warnings.warn(
```

=== Bagging ===

Accuracy: 0.74

	precision	recall	f1-score	support
-1	0.50	0.32	0.39	53
1	0.79	0.89	0.84	152
accuracy			0.74	205

macro avg	0.64	0.60	0.61	205
weighted avg	0.71	0.74	0.72	205

=== Boosting (AdaBoost) ===

Accuracy: 0.72

	precision	recall	f1-score	support
-1	0.43	0.28	0.34	53
1	0.78	0.87	0.82	152

accuracy			0.72	205
macro avg	0.60	0.58	0.58	205
weighted avg	0.69	0.72	0.70	205

=== Gradient Boosting ===

Accuracy: 0.71

	precision	recall	f1-score	support
-1	0.39	0.23	0.29	53
1	0.76	0.88	0.82	152

accuracy			0.71	205
macro avg	0.58	0.55	0.55	205
weighted avg	0.67	0.71	0.68	205